

# Part A: Mushroom Classification

## Prediction Task

The prediction task is to create a machine learning model to **predict a mushroom's edibility**, either edible or poisonous, based on the mushroom's attributes.

## Output Variable

The output variable is a **two-class** label - **edible** or **poisonous**. Edible (e) is defined the mushroom would cause no harm when consumed; whereas poisonous (p) refers to the mushroom will cause harmful effects when eaten.

# Data Profile

Overview

Warnings 47

Reproduction

## Dataset statistics

Number of variables	23
Number of observations	8124
Missing cells	2480
Missing cells (%)	1.3%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	1.4 MiB
Average record size in memory	184.0 B

## Variable types

Categorical	22
Boolean	1

# Warnings


Profile Reports shows **47 Warnings**

## Key points:

- All of the columns are categorical data
  - requires encoding
- “veil-type” contains only one class
  - drop column since there is no variance
- “stalk-root” has 30.5% missing data
  - either impute or drop the entire column
- high correlation between features
  - have to check for multicollinearity

Overview	Warnings 47	Reproduction
Warnings		
veil-type	has constant value "p"	Constant
class	is highly correlated with stalk-surface-above-ring and 10 other fields	High correlation
gill-attachment	is highly correlated with stalk-color-above-ring and 5 other fields	High correlation
stalk-surface-above-ring	is highly correlated with class and 11 other fields	High correlation
ring-type	is highly correlated with stalk-surface-above-ring and 11 other fields	High correlation
gill-spacing	is highly correlated with class and 6 other fields	High correlation
stalk-color-above-ring	is highly correlated with class and 15 other fields	High correlation
odor	is highly correlated with class and 15 other fields	High correlation
cap-surface	is highly correlated with stalk-root and 1 other fields	High correlation
population	is highly correlated with class and 14 other fields	High correlation
stalk-surface-below-ring	is highly correlated with class and 11 other fields	High correlation
ring-number	is highly correlated with ring-type and 7 other fields	High correlation
stalk-shape	is highly correlated with ring-type and 8 other fields	High correlation
habitat	is highly correlated with gill-spacing and 9 other fields	High correlation
stalk-color-below-ring	is highly correlated with class and 15 other fields	High correlation
stalk-root	is highly correlated with stalk-surface-above-ring and 14 other fields	High correlation
cap-shape	is highly correlated with population and 2 other fields	High correlation
spore-print-color	is highly correlated with class and 15 other fields	High correlation
cap-color	is highly correlated with ring-type and 11 other fields	High correlation
gill-size	is highly correlated with class and 6 other fields	High correlation
veil-color	is highly correlated with gill-attachment and 5 other fields	High correlation

# Multicollinearity Feature Selection



# Theil's U - Univariate Feature Selection

## Key Points to take down:

- These columns have low correlation with edibility, where  $U < 0.05$
- cap-shape, cap-surface, cap-color, gill-attachment, stalk-shape, veil-type, veil-color, ring-number

## Action Plan

- Drop all these columns show below

	cap-shape	cap-surface	cap-color	gill-attachment	stalk-shape	veil-type	veil-color	ring-number
class	0.048842	0.028617	0.036083	0.014178	0.007524	-1.603208e-16	0.023839	0.038489

# Pipeline/Feature Engineering



```
scoring = ["accuracy", "balanced_accuracy", "f1", 'roc_auc']

for name, model in models:
    pipeline = Pipeline(
        steps=[
            ("Imputation", SimpleImputer(strategy="most_frequent")),
            ("One-Hot Encoding", OneHotEncoder(drop="first", sparse=False)),
            (name, model)
        ]
    )
    cross_validate(pipeline, X_train, y_train, scoring=scoring)
```

# Model Selection


Most of the models I have preselected has 1.0 accuracy.

Since there is a dilemma in choosing learning algorithms, I decided to go with the simplest model with the best interpretability, Decision Tree Classifier.


	fit_time	score_time	test_accuracy	test_balanced_accuracy	test_f1	test_roc_auc
DecisionTreeClassifier	0.036299	0.015959	1.000000	1.000000	1.000000	1.000000
RandomForestClassifier	0.294108	0.050652	1.000000	1.000000	1.000000	1.000000
AdaBoostClassifier	0.549287	0.079787	1.000000	1.000000	1.000000	1.000000
GradientBoostingClassifier	1.052442	0.024467	1.000000	1.000000	1.000000	1.000000
Perceptron	0.043483	0.032352	1.000000	1.000000	1.000000	1.000000
Linear SVC	0.139350	0.035562	1.000000	1.000000	1.000000	1.000000
Polynomial SVC	0.359826	0.079387	1.000000	1.000000	1.000000	1.000000
ExtraTreesClassifier	0.372700	0.067840	1.000000	1.000000	1.000000	1.000000
MLPClassifier	1.763953	0.037759	1.000000	1.000000	1.000000	1.000000
CalibratedClassifierCV	0.203737	0.043484	1.000000	1.000000	1.000000	1.000000
SGDClassifier	0.048475	0.029321	1.000000	1.000000	1.000000	1.000000
KNeighborsClassifier	0.045479	0.701378	0.999472	0.999452	0.999450	1.000000
RBF SVC	0.421568	0.208354	0.999472	0.999452	0.999450	1.000000
RidgeClassifier	0.050712	0.033710	0.999120	0.999086	0.999084	1.000000
RidgeClassifierCV	0.109710	0.035540	0.999120	0.999086	0.999084	1.000000
LogisticRegression	0.112658	0.034307	0.999120	0.999086	0.999084	0.999987
GaussianNB	0.034707	0.025941	0.974147	0.975027	0.973790	0.996524
Sigmoid SVC	0.334529	0.079947	0.964649	0.964138	0.962766	0.985698

# Comparing to Baseline - Confusion Matrix

DecisionTreeClassifier K-Fold Cross Validation Confusion Matrix




DummyClassifier K-Fold Cross Validated Confusion Matrix






# Comparing to Baseline - Learning Curve

Learning Curve for DummyClassifier



Learning Curve for DecisionTreeClassifier




# Hyperparameter Tuning


```
[35] ▶ ML
filterwarnings('ignore')
# Create the parameter grid
params_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': np.arange(5, 15),
    'max_leaf_nodes': np.arange(10, 16)
}
# Creating a model based on the pipeline
grid_search = Pipeline(
    steps=[
        ('SimpleImputer', imp),
        ('OneHotEncoder', onehot),
        ('GridSearchCV', GridSearchCV(
            DecisionTreeClassifier(min_samples_split=2, min_samples_leaf=1),
            params_grid,
            cv=5,
            verbose=2,
            n_jobs=4,
            scoring='accuracy'
        ))
    ]
)
# Fitting Model
grid_search.fit(X_train, y_train)
print(grid_search.named_steps['GridSearchCV'].best_estimator_)
print(grid_search.named_steps['GridSearchCV'].best_params_)
print(grid_search.named_steps['GridSearchCV'].best_score_)

Fitting 5 folds for each of 120 candidates, totalling 600 fits
DecisionTreeClassifier(max_depth=7, max_leaf_nodes=14)
{'criterion': 'gini', 'max_depth': 7, 'max_leaf_nodes': 14}
1.0
```


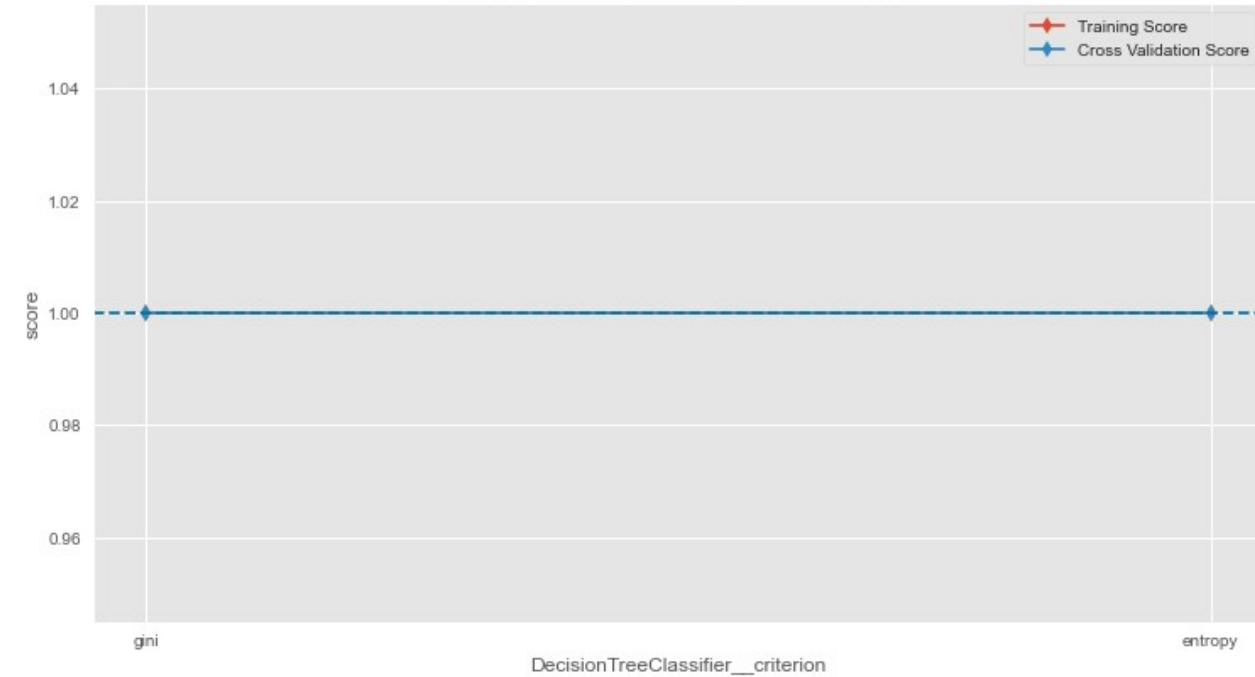
Validation Curve for DecisionTreeClassifier



Validation Curve for DecisionTreeClassifier




Validation Curve for DecisionTreeClassifier




# Model Evaluation - Learning Curve

Learning Curve for DecisionTreeClassifier Before Tuning



Learning Curve for DecisionTreeClassifier After Tuning



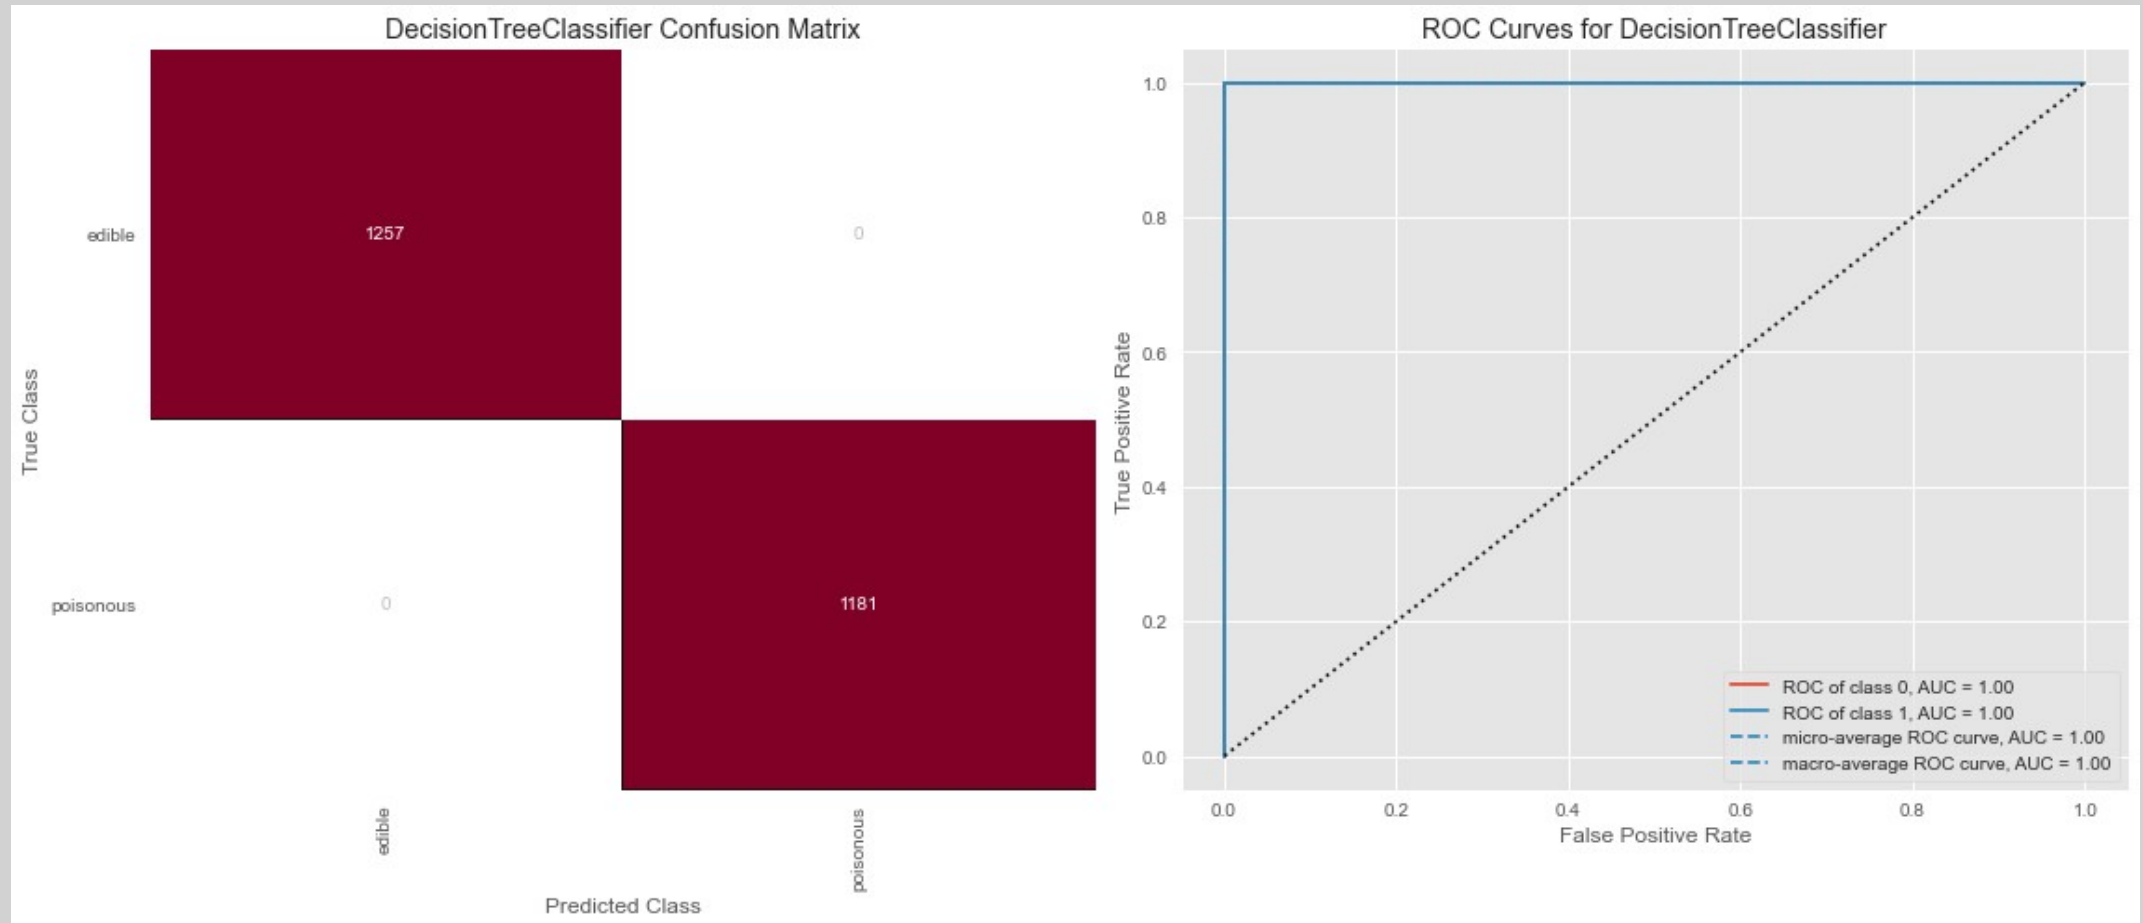
# Model Evaluation

	precision	recall	f1-score	support
0	1.000	1.000	1.000	1257
1	1.000	1.000	1.000	1181
accuracy			1.000	2438
macro avg	1.000	1.000	1.000	2438
weighted avg	1.000	1.000	1.000	2438

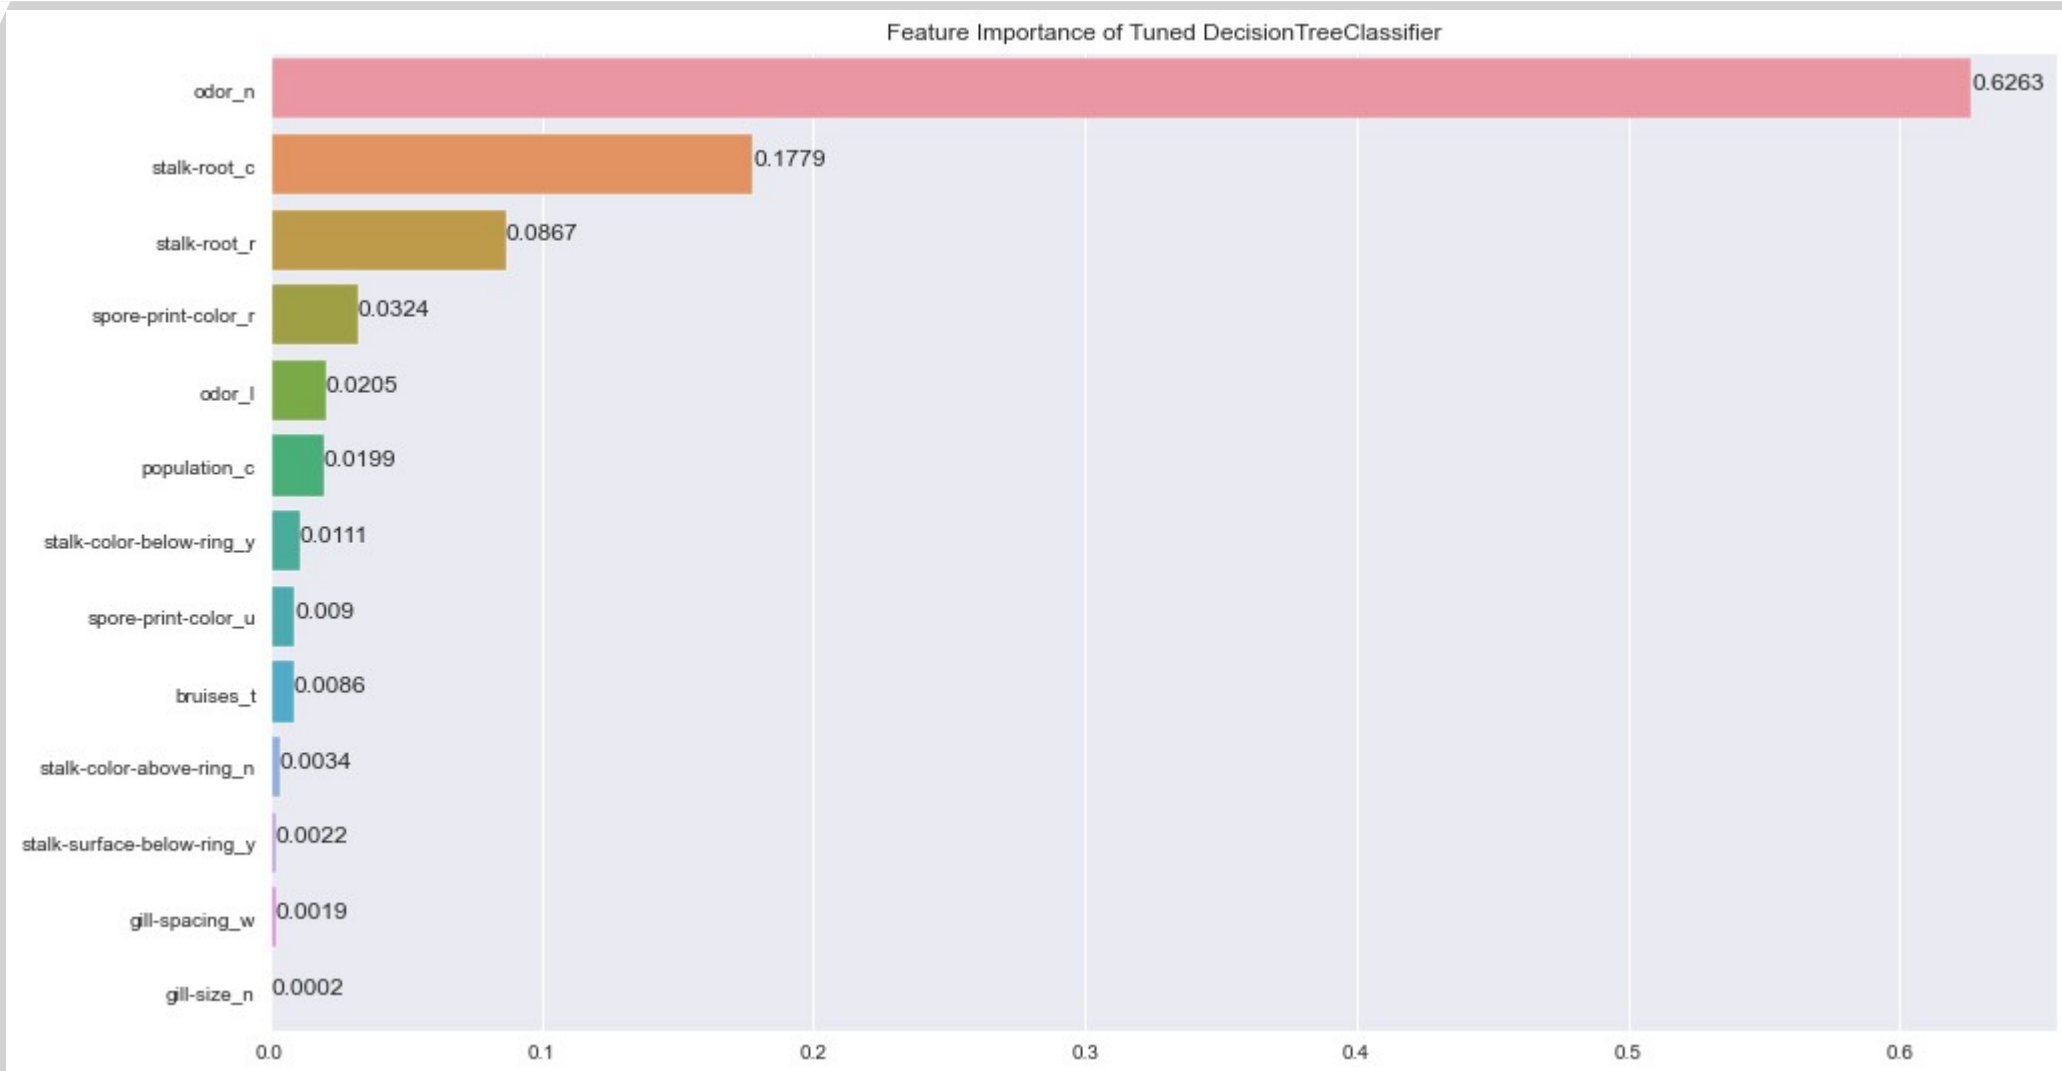
Cross Validation Scores: [1. 1. 1. 1. 1.]

Mean Cross Validation Scores: 1.0

# Model Evaluation



# Model Evaluation - Feature Importance



# Part B: Kings County House Price Prediction

## **Prediction Task**

The prediction task is to create a predictive regression model to predict the house prices based on the house sales' attributes given.

## **Output Variable**

The output variable is `price`, referring to the price of the house sale. The variable is a numerical-continuous variable, as such the prediction task requires a regression model.



# Data Profile

Overview

Warnings 55

Reproduction

## Dataset statistics

Number of variables	21
Number of observations	21613
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	3.5 MiB
Average record size in memory	168.0 B

## Variable types

Numeric	17
DateTime	1
Categorical	3

# Warnings


Profile Reports shows **55 Warnings**

## Key points:

- high correlation between features
  - have to check for multicollinearity
- 'sqft\_basement' has 60.7% zeros
- 'yr\_renovated' has 95.8% zeros

sqft_basement is highly correlated with sqft_above and 3 other fields	High correlation
sqft_above is highly correlated with sqft_living15 and 6 other fields	High correlation
zipcode is highly correlated with lat and 2 other fields	High correlation
sqft_living is highly correlated with sqft_living15 and 6 other fields	High correlation
long is highly correlated with zipcode and 1 other fields	High correlation
grade is highly correlated with sqft_living15 and 4 other fields	High correlation
bathrooms is highly correlated with sqft_living15 and 7 other fields	High correlation
yr_built is highly correlated with condition and 4 other fields	High correlation
bedrooms is highly correlated with sqft_above and 2 other fields	High correlation
sqft_lot15 is highly correlated with sqft_lot	High correlation
price is highly correlated with sqft_living15 and 5 other fields	High correlation
view is highly correlated with waterfront	High correlation
waterfront is highly correlated with view	High correlation
sqft_basement has 13126 (60.7%) zeros	Zeros
yr_renovated has 20699 (95.8%) zeros	Zeros

# Pearson Correlation Matrix




# House Age


Age referring to the time difference between date and yr\_renovated/yr\_built

	date	yr_renovated	yr_built	age
0	2014-10-13	0	1955	59.0
1	2014-12-09	1991	1951	23.0
2	2015-02-25	0	1933	82.0
3	2014-12-09	0	1965	49.0
4	2015-02-18	0	1987	28.0
...	...	...	...	...
21608	2014-05-21	0	2009	5.0
21609	2015-02-23	0	2014	1.0
21610	2014-06-23	0	2009	5.0
21611	2015-01-16	0	2004	11.0
21612	2014-10-15	0	2008	6.0


21613 rows x 4 columns




# Log Transformation




Before Normalisation




After Normalisation



House Prices before Transformation



House Prices after Log Tranformation



# Model Selection

```
for name, models in models:  
    pipeline = Pipeline(  
        steps=[  
            ('LogTransform', LogTransform),  
            ('Normalisation', scaler),  
            (name, TransformedTargetRegressor(regressor=model, func=np.log1p, inverse_func=np.expm1))  
        ]  
    )  
    cross_validate(pipeline, X_train, y_train, cv=5, scoring=['RMSE', 'MSE', 'MAE', 'MAPE', 'R2'])
```





# Model Selection

	fit_time	score_time	test_rmse	train_rmse	test_mse	train_mse	test_mae	train_mae	test_mape	train_mape	test_r2	train_r2
ExtraTreesRegressor	7.135958	0.200463	-133390.500654	-3068.189155	-17936571347.077385	-9513161.568014	-69219.727082	-123.724826	-0.126745	-0.000425	0.869151	0.999930
RandomForestRegressor	12.393565	0.149601	-135627.651945	-55290.737436	-18525893661.364635	-3058882980.024376	-70347.560765	-26816.512434	-0.128178	-0.047119	0.864615	0.977534
GradientBoostingRegressor	3.347374	0.010773	-139285.123516	-128246.897591	-19486473186.692482	-16455351352.906652	-75522.579562	-71986.147128	-0.137311	-0.131361	0.856904	0.879157
AdaBoostRegressor	1.306047	0.024340	-140606.259676	-64731.613580	-19908085622.554283	-4203027839.285998	-74223.345925	-30492.814420	-0.135643	-0.053883	0.854434	0.969156
BaggingRegressor	1.309004	0.028725	-140874.161656	-63135.779721	-19951648787.887047	-3996438359.867338	-74411.359193	-30475.733821	-0.136075	-0.053843	0.853591	0.970632
SVR	13.303492	4.843579	-145188.994993	-131623.979113	-21304507975.040260	-17330699587.146801	-76776.802559	-72849.505850	-0.139376	-0.134285	0.844945	0.872686
MLPRegressor	15.115033	0.012760	-149213.947738	-141366.393039	-22478985541.641655	-20019688178.613991	-79865.868099	-78126.303800	-0.142699	-0.140868	0.835492	0.852696
KNeighborsRegressor	0.189692	0.985876	-168784.280528	-140807.627813	-28818222820.154930	-19848924340.629433	-84608.095344	-68939.526920	-0.150401	-0.121234	0.790136	0.854273
DecisionTreeRegressor	0.213815	0.008178	-188055.943134	-3068.189155	-35613540755.972511	-9513161.568014	-100716.780634	-123.724826	-0.185233	-0.000425	0.739723	0.999930
LinearRegression	0.624596	0.075797	-188994.705372	-188944.478452	-35841165312.482529	-35707235447.812317	-110814.325837	-110672.289058	-0.201534	-0.201302	0.736630	0.737656
Ridge	0.368858	0.006184	-189350.393436	-189301.885531	-35977847205.701393	-35842467227.082199	-110861.627598	-110719.515040	-0.201544	-0.201316	0.735637	0.736663
SGDRegressor	0.138829	0.007580	-210493.017840	-210375.920353	-44405094031.160751	-44267019889.859184	-125276.401612	-125086.255084	-0.221620	-0.221336	0.672652	0.674762
ElasticNet	0.051257	0.007977	-376119.759146	-376659.424464	-141907449864.942780	-141897913482.805664	-221876.211038	-221875.234014	-0.438628	-0.438620	-0.042641	-0.042515
Lasso	0.021941	0.007579	-376119.759146	-376659.424464	-141907449864.942780	-141897913482.805664	-221876.211038	-221875.234014	-0.438628	-0.438620	-0.042641	-0.042515

# Comparing with Baseline

	fit_time	score_time	test_mae	test_mape	test_mse	test_r2	test_rmse
GradientBoostingRegressor	3.382798	0.009982	-75506.759853	-0.137306	-1.943598e+10	0.857226	-139112.218952
LinearRegression	0.021336	0.007188	-110814.325837	-0.201534	-3.584117e+10	0.736630	-188994.705372



# Hyperparameter Tuning - HalvingGridSearchCV

```
grid = {
    'regressor__n_estimators': np.arange(100, 2101, 200),
    'regressor__max_depth': [3, 5, 10],
    'regressor__learning_rate': [.001, .01, .1],
    'regressor__subsample': [.5, .75, 1]
}

model_tuning = Pipeline(
    steps=[
        ('LogTransform', LogTransform()),
        ('Normalisation', scaler),
        ('GridSearch', HalvingGridSearchCV(
            TransformedTargetRegressor(
                regressor=GradientBoostingRegressor(),
                func=np.log1p,
                inverse_func=np.expm1
            ),
            grid,
            scoring='neg_root_mean_squared_error',
            n_jobs=4,
            verbose=1,
            cv=5,
            aggressive_elimination=True,
            factor=5
        ))
    ]
)

model_tuning.fit(X_train, y_train)
print('Finish Tuning')
```

# Hyperparameter Tuning

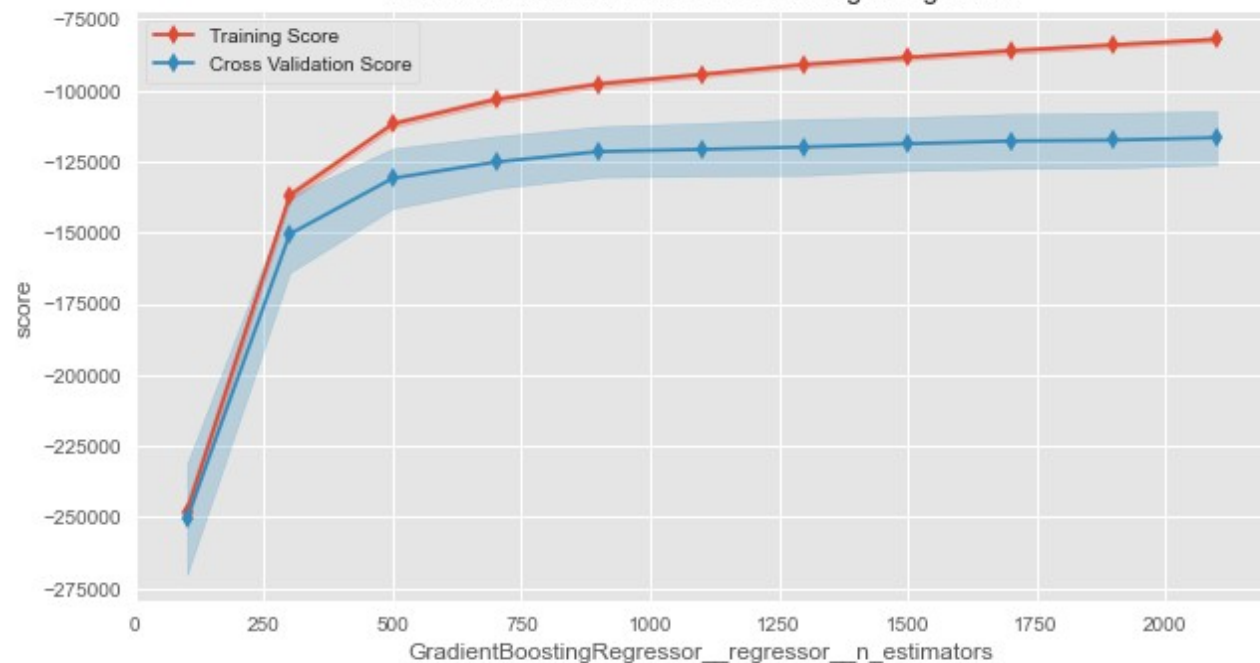
## - HalvingGridSearchCV

▶ M4


```
print(model_tuning.named_steps['GridSearch'].best_estimator_)  
print(model_tuning.named_steps['GridSearch'].best_params_)  
print(model_tuning.named_steps['GridSearch'].best_score_)
```

```
TransformedTargetRegressor(func=<ufunc 'log1p'>, inverse_func=<ufunc 'expm1'>,  
                           regressor=GradientBoostingRegressor(learning_rate=0.01,  
                                                                max_depth=5,  
                                                                n_estimators=1700,  
                                                                subsample=0.5))  
{'regressor__learning_rate': 0.01, 'regressor__max_depth': 5, 'regressor__n_estimators': 1700, 'regressor__subsample': 0.5}  
-117994.0405785848
```


Validation Curve for TransformedTargetRegressor



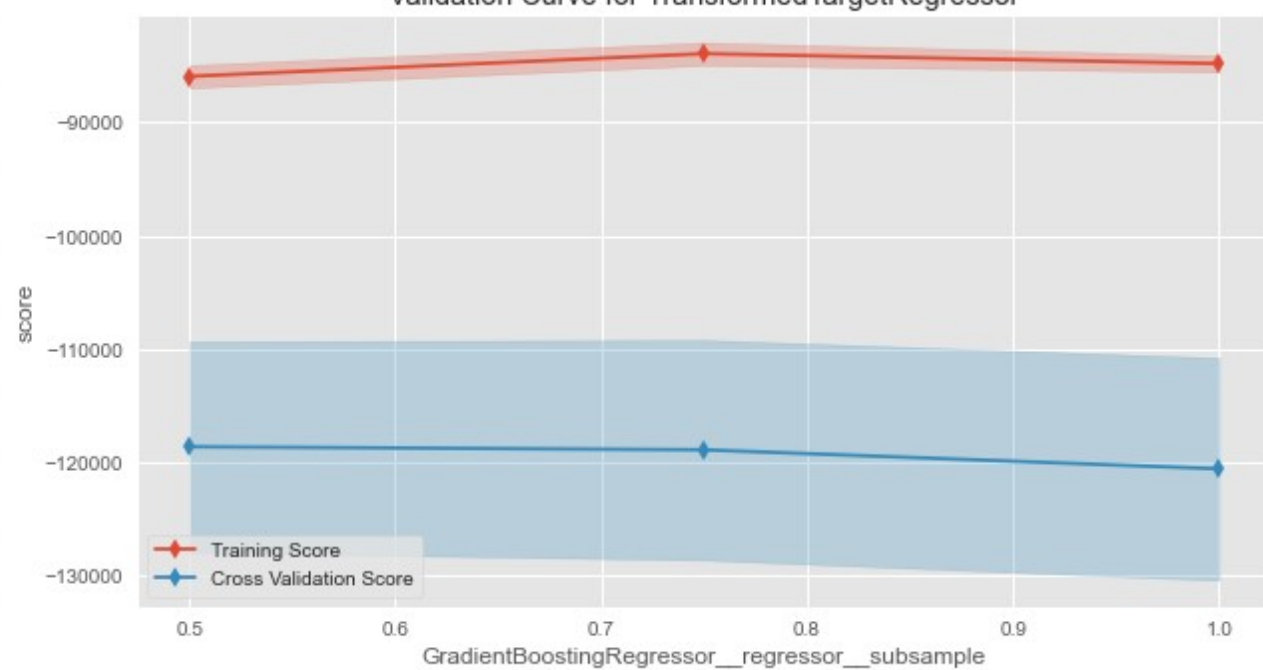
Validation Curve for TransformedTargetRegressor




Validation Curve for TransformedTargetRegressor




Validation Curve for TransformedTargetRegressor



### Learning Curve for GradientBoostingRegressor without Hyperparameters





### Learning Curve for GradientBoostingRegressor with Hyperparameters





# Generating Predictions

	RMSE	MSE	MAE	MAPE	R2
GradientBoostingRegressor	112580.789489	1.267443e+10	64669.690289	0.121199	0.982068









# Feature Importance

